

MPI 2

1)MPI_Gather() & MPI_Scatter()

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <mpi.h>
#include <assert.h>

// Creates an array of random numbers. Each number has a value from 0
- 1
float *create_rand_nums(int num_elements) {
    float *rand_nums = (float *)malloc(sizeof(float) * num_elements);
    assert(rand_nums != NULL);
    int i;
    for (i = 0; i < num_elements; i++) {
        rand_nums[i] = (rand() / (float)RAND_MAX);
    }
    return rand_nums;
}

// Computes the average of an array of numbers
float compute_avg(float *array, int num_elements) {
    float sum = 0.f;
    int i;
    for (i = 0; i < num_elements; i++) {
        sum += array[i];
    }
    return sum / num_elements;
}

int main(int argc, char** argv) {
    if (argc != 2) {
        fprintf(stderr, "Usage: avg num_elements_per_proc\n");
        exit(1);
    }

    int num_elements_per_proc = atoi(argv[1]);
    // Seed the random number generator to get different results each
time
    srand(time(NULL));
```

```

MPI_Init(NULL, NULL);

int world_rank;
MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
int world_size;
MPI_Comm_size(MPI_COMM_WORLD, &world_size);

// Create a random array of elements on the root process. Its total
// size will be the number of elements per process times the number
// of processes
float *rand_nums = NULL;
if (world_rank == 0) {
    rand_nums = create_rand_nums(num_elements_per_proc *
world_size);
}

// For each process, create a buffer that will hold a subset of the
entire
// array
float *sub_rand_nums = (float *)malloc(sizeof(float) *
num_elements_per_proc);
assert(sub_rand_nums != NULL);

// Scatter the random numbers from the root process to all
processes in
// the MPI world
MPI_Scatter(rand_nums, num_elements_per_proc, MPI_FLOAT,
sub_rand_nums,
            num_elements_per_proc, MPI_FLOAT, 0, MPI_COMM_WORLD);

// Compute the average of your subset
float sub_avg = compute_avg(sub_rand_nums, num_elements_per_proc);

// Gather all partial averages down to the root process
float *sub_avgs = NULL;
if (world_rank == 0) {
    sub_avgs = (float *)malloc(sizeof(float) * world_size);
    assert(sub_avgs != NULL);
}
MPI_Gather(&sub_avg, 1, MPI_FLOAT, sub_avgs, 1, MPI_FLOAT, 0,
MPI_COMM_WORLD);

```

```

    // Now that we have all of the partial averages on the root,
compute the
    // total average of all numbers. Since we are assuming each process
computed
    // an average across an equal amount of elements, this computation
will
    // produce the correct answer.
    if (world_rank == 0) {
        float avg = compute_avg(sub_avgs, world_size);
        printf("Avg of all elements is %f\n", avg);
        // Compute the average across the original data for comparison
        float original_data_avg =
            compute_avg(rand_nums, num_elements_per_proc * world_size);
        printf("Avg computed across original data is %f\n",
original_data_avg);
    }

    // Clean up
    if (world_rank == 0) {
        free(rand_nums);
        free(sub_avgs);
    }
    free(sub_rand_nums);

    MPI_Barrier(MPI_COMM_WORLD);
    MPI_Finalize();
}

```

Output:

```

mnit@mnit-OptiPlex-5040:~$ mpirun -np 4 ./scaga 3
Avg of all elements is 0.519614
Avg computed across original data is 0.519614
mnit@mnit-OptiPlex-5040:~$ mpirun -np 4 ./scaga 5
Avg of all elements is 0.435156
Avg computed across original data is 0.435156
mnit@mnit-OptiPlex-5040:~$ gedit scaga.c

```

2) MPI_Get_count()

```

#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

```

```

int main(int argc, char** argv) {
    MPI_Init(NULL, NULL);

    int world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);
    if (world_size != 2) {
        fprintf(stderr, "Must use two processes for this example\n");
        MPI_Abort(MPI_COMM_WORLD, 1);
    }
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);

    int number_amount;
    if (world_rank == 0) {
        const int MAX_NUMBERS = 100;
        int numbers[MAX_NUMBERS];
        // Pick a random amount of integers to send to process one
        srand(time(NULL));
        number_amount = (rand() / (float)RAND_MAX) * MAX_NUMBERS;
        // Send the amount of integers to process one
        MPI_Send(numbers, number_amount, MPI_INT, 1, 0, MPI_COMM_WORLD);
        printf("0 sent %d numbers to 1\n", number_amount);
    } else if (world_rank == 1) {
        MPI_Status status;
        // Probe for an incoming message from process zero
        MPI_Probe(0, 0, MPI_COMM_WORLD, &status);
        // When probe returns, the status object has the size and other
        // attributes of the incoming message. Get the size of the
        message.
        MPI_Get_count(&status, MPI_INT, &number_amount);
        // Allocate a buffer just big enough to hold the incoming numbers
        int* number_buf = (int*)malloc(sizeof(int) * number_amount);
        // Now receive the message with the allocated buffer
        MPI_Recv(number_buf, number_amount, MPI_INT, 0, 0,
MPI_COMM_WORLD,
                MPI_STATUS_IGNORE);
        printf("1 dynamically received %d numbers from 0.\n",
                number_amount);
        free(number_buf);
    }
    MPI_Finalize();
}

```

Output:

```
mnit@mnit-OptiPlex-5040:~$ mpirun -np 2 ./count
0 sent 73 numbers to 1
1 dynamically received 73 numbers from 0.
```

3) MPI_Issend() & MPI_Irecv()

```
#include "mpi.h"
#include <stdio.h>
int main(int argc, char *argv[])
{
    int myid, numprocs, left, right, i;
    int buffer[10], buffer2[10];
    MPI_Request request, request2;
    MPI_Status status;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);
    right = (myid + 1) % numprocs;
    left = myid - 1;
    if (left < 0)
        left = numprocs - 1;
    MPI_Irecv(buffer, 10, MPI_INT, left, 123, MPI_COMM_WORLD,
&request);
    MPI_Issend(buffer2, 10, MPI_INT, right, 123, MPI_COMM_WORLD,
&request2);
    for(i=0; i<10; i++)
        printf("Value is:%d\n", buffer[i]);
    MPI_Wait(&request, &status);
    MPI_Wait(&request2, &status);
}
```

```
    MPI_Finalize();  
    return 0;  
}
```

Output:

```
mnit@mnit-OptiPlex-5040:~/n6$ mpirun -np 4 5  
Value is:1254359544  
Value is:32574  
Value is:1266391760  
Value is:32574  
Value is:0  
Value is:0  
Value is:-1276493800  
Value is:32764  
Value is:1  
Value is:0  
Value is:2105172472  
Value is:32753  
Value is:2117204688  
Value is:32753  
Value is:0  
Value is:0  
Value is:-731666664  
Value is:32766  
Value is:1  
Value is:0  
Value is:1408020984  
Value is:32661  
Value is:1420053200  
Value is:32661  
Value is:0  
Value is:0  
Value is:2143762648  
Value is:32766  
Value is:1  
Value is:0  
Value is:1605657080  
Value is:32563  
Value is:1617689296  
Value is:32563  
Value is:0  
Value is:0
```

5) MPI_Test()

```
#include <stdio.h>
#include <mpi.h>
#include <sys/time.h>

int main(int argc, char *argv[])
{

    int rank, size;
    MPI_Status status;

    /* Init */
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    if (rank != 0) { // Slaves
        int buf;

        if (rank == 1) {
            buf = 1;
            MPI_Send(&buf, 1, MPI_INT, 0, 0, MPI_COMM_WORLD);
        }
        if (rank == 2) {
            buf = 2;
            MPI_Send(&buf, 1, MPI_INT, 0, 0, MPI_COMM_WORLD);
        }

    }
    else { // Master
        int sum = 0;
        int flag = -1, res;
        MPI_Request request;
        MPI_Status status;
        while (1) {
            if(flag != 0)
            {
                MPI_Irecv(&res, 1, MPI_INT, MPI_ANY_SOURCE, MPI_ANY_TAG,
MPI_COMM_WORLD, &request);
```

```

    flag = 0;
}
MPI_Test(&request, &flag, &status);

if (flag != 0) {
    printf("recv : %d, slave : %d\n", res, status.MPI_SOURCE);
    if (status.MPI_SOURCE != -1)
        sum += res;
    flag = -1;
}

if (sum == 3)
    break;
}

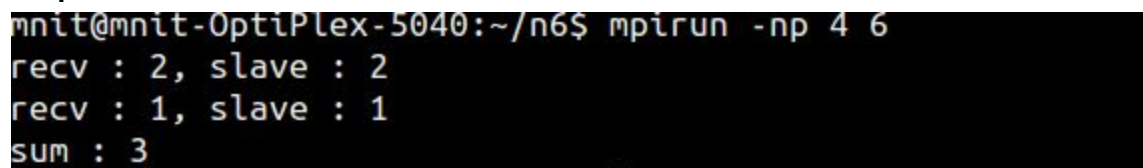
printf("sum : %d\n", sum);
}

MPI_Finalize();
return 0;

}

```

Output:



```

mnit@mnit-OptiPlex-5040:~/n6$ mpirun -np 4 6
recv : 2, slave : 2
recv : 1, slave : 1
sum : 3

```

6) MPI_Allgather()

```

#include <stdio.h>
#include <stdlib.h>
#include "mpi.h"

struct mystruct{
    int sendarray[4];
    int a;
    char array2[4];
};

typedef struct mystruct struct_t;

```



```

int main (int argc, char ** argv)
{
    int rank, size;
    struct_t fd;
    struct_t recv;
    int i, j;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    // init
    for (i=0;i<4;i++){
        fd.sendarray[i] = 0;
        fd.array2[i] = 0;

        recv.sendarray[i] =999;
        recv.array2[i]  = 99;
    }
    recv.a =999;

    // put some stuff in your array
    fd.sendarray[rank] = rank*10;
    fd.array2[rank] = (char)(rank*20);
    fd.a = rank;
    printf("My rank is %d, fd.sendarray[%d]  is %d\n", rank, rank,
fd.sendarray[rank]);

    // gather data from all now.. send the int:
    MPI_Allgather (&(fd.sendarray[rank]), 1, MPI_INT,
recv.sendarray, 1, MPI_INT,  MPI_COMM_WORLD);
    // then the char
    MPI_Allgather (&(fd.array2[rank]), 1, MPI_CHAR, recv.array2,
1, MPI_CHAR, MPI_COMM_WORLD);

    // check if correct data has been received
    if (rank == 0) {
        printf("Received:\n");
        printf("---\n");
        printf("int array:  ");
        for (j=0; j<4; j++) printf("%3d ", recv.sendarray[j]);
        printf("\nchar array: ");
        for (j=0; j<4; j++) printf("%3d ", (int)(recv.array2[j]));
        printf("\n");
    }
}

```

```

    }
    MPI_Finalize();

    return 0;
}

```

Output:

```

mnit@mnit-OptiPlex-5040:~/n6$ mpirun -np 4 7
My rank is 0, fd.sendarray[0] is 0
My rank is 1, fd.sendarray[1] is 10
My rank is 2, fd.sendarray[2] is 20
My rank is 3, fd.sendarray[3] is 30
Received:
---
int array:    0  10  20  30
char array:   0  20  40  60

```

7) MPI_Alltoall()

```

#include <stdio.h>
#include "mpi.h"

int main( int argc, char **argv )
{
    int send[4], recv[3];
    int rank, size, k;

    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &rank );
    MPI_Comm_size( MPI_COMM_WORLD, &size );

    if (size != 4) {
        printf("Error!:# of processors must be equal to 4");
        printf("Programm aborting....");
        MPI_Abort(MPI_COMM_WORLD, 1);
    }
    for (k=0;k<size;k++)
        send[k] = (k+1) + rank*size;

    printf("%d : send = %d %d %d %d\n", rank, send[0], send[1],
send[2], send[3]);
}

```

```
        MPI_Alltoall(&send, 1, MPI_FLOAT, &recv, 1, MPI_INT,
MPI_COMM_WORLD);

        printf("%d : recv = %d %d %d %d\n", rank, recv[0], recv[1],
recv[2], recv[3]);

        MPI_Finalize();
        return 0;
}
```

Output:

```
mnit@mnit-OptiPlex-5040:~/n6$ mpirun -np 4 8
0 : send = 1 2 3 4
1 : send = 5 6 7 8
2 : send = 9 10 11 12
3 : send = 13 14 15 16
1 : recv = 2 6 10 14
2 : recv = 3 7 11 15
3 : recv = 4 8 12 16
0 : recv = 1 5 9 13
```